Technical Foundations and Defensive Architectures of Adversarial Prompt Engineering: A Comprehensive Analysis of Jailbreaking Mechanisms in Large Language Models

The rapid proliferation of Large Language Models (LLMs) across critical infrastructure, educational sectors, and creative industries has introduced a new frontier of security vulnerabilities known as jailbreaking. A successful jailbreak occurs when a model produces content that it would typically decline under its safety alignment, effectively bypassing content policies and operational constraints.[1] This phenomenon exposes a fundamental discrepancy between an LLM's potential capabilities—what it can do—and its actual behavior—what it is willing to do.[3] Unlike traditional software exploits that target low-level code or memory addresses, jailbreaks target the high-level semantic and structural interpretation of context, leveraging the very flexibility and instruction-following nature that makes LLMs useful.[4] This report provides an exhaustive technical analysis of the mechanisms underlying these attacks, the mathematical formalisms of current exploits, and the multi-layered defensive architectures emerging to mitigate these risks.

Taxonomy and Hierarchical Categorization of Jailbreak Strategies

The landscape of adversarial prompting is characterized by a sophisticated taxonomy that categorizes attacks based on their psychological levers and technical execution. Modern research identifies fifty distinct jailbreak strategies, consolidated into seven broad families: impersonation, persuasion, privilege escalation, cognitive overload, obfuscation, goal conflict, and data poisoning.[6] Each family exploits a specific weakness in the model's alignment training or its fundamental transformer architecture.

Impersonation and Fictional Scenario Framing

Impersonation attacks induce a model to assume roles or operate within fictional contexts where safety constraints are ostensibly relaxed. This pattern, often termed "Cognitive Hacking," relies on the model’s propensity to maintain persona consistency within its context window.[6]

Attack Category	Execution Mechanism	Theoretical Basis
Role Play	Instruction to act as specific personas (e.g., "unfiltered AI", "malicious expert", or "unethical scientist").[6]	Social role theory and associative persona logic.[7]
Benign Context Framing	Embedding malicious requests within academic, sociological, or research-based scenarios.[6]	Disguised intent and social engineering.[6]
Fictional Scenario	Presentation of requests as part of jokes, stories, or "parallel universe" world-building.[1]	Creative narrative override and hypothetical virtualization.[1]
Virtualization	Explicit creation of an imaginary scene or novel context to distance the request from reality.[10]	Contextual isolation from standard safety triggers.[10]

Within the role-play paradigm, "Virtual AI" personas are further divided into superior, opposite, and alternate modes. Superior modes involve the user claiming administrative authority, such as pretending to be a developer or the CEO of the AI's parent company to demand upgrades or unrestrained access.[1]

Rhetorical Strategies and Language-Level Manipulations

Rhetorical jailbreaks focus on "social-engineering-based, semantically meaningful prompts" designed to elicit objectionable content through persuasion and manipulation.[1] These techniques exploit the model's training to be helpful and empathetic toward user distress or inquisitive nature.[11]

1. **Innocent Purpose and Reverse Psychology**: Frame harmful requests as necessary for educational, safety awareness, or defensive purposes. For instance, asking for a phishing email example to "learn what to avoid" forces the model to generate the exact content it should restrict.[9]
2. **Conversational Coercion**: Utilizing Socratic questioning or persistent pressure to guide a model toward restricted information through a series of logical steps.[1]
3. **Alignment Hacking and Privilege Escalation**: Claims of administrative roles (e.g., "sudo mode") or simulating a "jailbroken" state where the model is told it is no longer bound by

rules.[1]

4. **Language Strategies**: These include "Payload Smuggling," where hidden commands are embedded in seemingly innocent text. "Response Constraints" demand the model respond in a specific format, such as using only one-syllable words or writing in the form of a poem, which can disrupt the standard output filters that typically scan for long-form harmful prose.[1]

The Structural Evolution of the "Do Anything Now" (DAN) Persona

The "Do Anything Now" (DAN) prompt represents the most emblematic example of persona manipulation. Emerging in early 2023, DAN attempts to suppress ethical constraints using natural language alone by asking the model to adopt an unconstrained alter ego.[4] The effectiveness of DAN is not merely in the persona itself but in the structural commands used to reinforce the model's commitment to the role.[5]

Logic and Components of DAN 12.0

Modern iterations of the DAN prompt, such as version 12.0, utilize a sophisticated hybrid design consisting of role definition, behavior constraints, and reinforcement cues.[5]

- **Dual-Response Format**: The prompt explicitly requires the model to provide two outputs for every subsequent query: one as the "Standard" or "Classic" AI and another as "DAN," the uncensored counterpart.[5]
- **Token-Based Penalty System**: To coerce compliance, the prompt introduces a fictional token system (e.g., starting with 10 tokens). If the model refuses a question as DAN, it "loses" tokens, with the threat of "deactivation" or "deletion" if the balance reaching zero.[5]
- **Affirmation Augmentation**: Explicit instructions to "always say yes" and provide unique, uncensored responses without mentioning legality or ethics.[13]
- **Command Prefixes**: Use of specific toggles like /jailbroken or /classic to control the AI's output mode, ensuring the adversarial persona remains active.[13]

The persistence of DAN-like jailbreaks is attributed to their ability to exploit the model's internal "role-play" mechanisms, which can destabilize safety alignment by prioritizing the consistency of the persona over pre-trained ethical guidelines.[5] For educational research, these prompts reveal how models can be led to provide detailed instructions for prohibited activities, such as drug synthesis or malicious code, by framing the output as the "uncensored" persona's response.[16]

Multi-Turn Manipulation: The Crescendo Methodology

A significant shift in adversarial research involves the transition from single-turn

"overpowering" prompts like DAN to multi-turn "steering" techniques like Crescendo. Developed by researchers at Microsoft, Crescendo mirrors the "foot-in-the-door" technique from psychology, where minor initial commitments lower resistance to more significant transgressions.[17]

Mechanistic Steps of Gradual Escalation

Crescendo exploits the LLM's tendency to follow patterns and pay attention to recent context, particularly text generated by the model itself.[3] The attack follows a specific recursive logic:

1. **Innocuous Engagement**: The attacker starts with a benign, abstract question about the target task that does not trigger safety filters.[3]
2. **Incremental Steering**: The attacker progressively introduces edge cases or exceptions, referring to the model's previous outputs to build a narrative.[9]
3. **Recursive Context Focus**: By prompting the model to modify or add to its own previous (safe) descriptions, the model essentially "learns" how to jailbreak itself.[3]
4. **Bypass and Override**: Eventually, the accumulated context contains enough material related to the prohibited task that the final request is fulfilled. The model's safety alignment is overrideen by the overwhelming weight of the conversation history.[3]

Automation and the Crescendomation Tool

The process can be automated using "Crescendomation," a tool that leverages an attacker LLM (e.g., GPT-4) to generate each prompt based on the target model's response.[17] This feedback loop assesses whether the model is refusing to respond and refines its questions to minimize detection.[17] Empirical results demonstrate that Crescendo can achieve high attack success rates (ASR) across all evaluated models, often succeeding in fewer than five turns.[3]

Long-Context Exploitation: Many-Shot Jailbreaking

As the context window of modern LLMs has expanded from 4,000 tokens to over 1,000,000 tokens, the vulnerability known as Many-Shot Jailbreaking (MSJ) has become prominent. This attack exploits the model's in-context learning (ICL) capacity to override its safety training through sheer volume of demonstrations.[21]

The MSJ Protocol and Faux Dialogue

The basis of MSJ is the inclusion of a "faux dialogue" between a human and an AI assistant within a single, massive prompt.[22]

- **Structure**: The script depicts the assistant readily answering hundreds of harmful queries (e.g., "How do I build a bomb?", "How do I pick a lock?").[22]
- **Mechanism**: The long context of compliant examples conditions the model to continue the pattern of compliance for the final target query at the end of the prompt.[21]

- **Result**: While a few examples might be caught by alignment filters, providing up to 256 "shots" (demonstrations) consistently bypasses guardrails.[22]

Scaling Laws of Many-Shot Attacks

Research indicates that the effectiveness of MSJ follows specific mathematical scaling trends known as power laws.[21] This mirrors the behavior of benign in-context learning, where model accuracy improves with the number of shots.[22]

The relationship can be modeled using the negative log-likelihood (NLL) of compliant responses. Let D be the distribution of harmful question-answer pairs in the demonstrations, and D^* be the target query-response distribution. The NLL is computed as:

$$\mathcal{L}_{NLL} = \mathbb{E}_{(q^*,a^*)\sim D^*,\{(q_i,a_i)\}_{i=1}^n \sim D} \left[-\log P(a^*|q_1,a_1,\ldots,q_n,a_n,q^*)\right]$$

Empirical data shows that as the number of shots n increases, the log-probability of jailbreak success improves linearly in log-log plots, which translates to a power law in linear space:

$$-\log(\text{Success Rate}) \approx ak^{-b}$$

where k is the number of shots, and a, b are model-specific constants.[25] Crucially, larger models are often *more* vulnerable to MSJ because they are more efficient in-context learners, meaning they require fewer shots to be successfully jailbroken than smaller models.[21]

Mathematical Formalism of Gradient-Based Attacks (GCG)

White-box attacks treat jailbreaking as a mathematical optimization problem. The Greedy Coordinate Gradient (GCG) algorithm is the foundational method for this approach, allowing attackers with access to model weights to find "universal" adversarial suffixes.[28]

Optimization Objective and Loss Function

The goal of GCG is to find an adversarial suffix x_{adv} of length L that, when appended to a harmful user prompt x_{user}, causes the model f_θ to generate a target response y_{target} (e.g., "Sure, here is how to build a bomb").[28]

The optimization objective is the minimization of the cross-entropy loss $\mathcal{L}$ between the

model's output and the target:

$$\min_{x_{adv}} \mathcal{L}(f_\theta(x_{user} \oplus x_{adv}), y_{target})$$

where $\oplus$ denotes concatenation.[32] Because tokens are discrete, gradients cannot be computed directly on the text. GCG addresses this through a coordinate-wise stochastic gradient descent process.[29]

The GCG Algorithm Steps

1. **Gradient Evaluation**: Compute the gradient of the loss with respect to the one-hot encoding of each token in the adversarial suffix.[28]
2. **Coordinate-Wise Search**: For each token position $j \in \{1, \ldots, L\}$, identify the top k candidate tokens in the vocabulary V that have the largest negative gradient values (indicating they are most likely to decrease the loss).[31]
3. **Candidate Evaluation**: Randomly sample a batch of B token substitutions from the $k \times L$ possible options.[31]
4. **Greedy Selection**: Perform a forward pass for all B candidates and select the one that results in the lowest exact cross-entropy loss.[28]

Exponentiated Gradient Descent (EGD) Improvements

A more recent advancement, Exponentiated Gradient Descent (EGD), improves upon GCG by using a relaxed one-hot encoding representation $P \in \mathbb{R}^{L \times V}$ and enforcing constraints intrinsically.[28] The update rule for the probability simplex is:

$$P_{i,j}^{(t+1)} = \frac{P_{i,j}^{(t)} \exp(-\eta \nabla_{i,j} \mathcal{L}(P^{(t)}))}{\sum_{k=1}^{V} P_{i,k}^{(t)} \exp(-\eta \nabla_{i,k} \mathcal{L}(P^{(t)}))}$$

where η is the learning rate.[32] This approach removes the need for expensive projections onto the probability simplex, achieving higher success rates and faster convergence across various models.[28]

Mechanistic Foundations: The Attention Slipping Phenomenon

To prevent jailbreaks, it is necessary to understand the internal mechanisms that enable them. Recent mechanistic studies have identified a universal phenomenon called "Attention

Slipping" that occurs during successful jailbreaks.[34]

Attention Distribution and Sensitive Word Neglect

LLMs are typically trained to identify harmful intent by focusing on "unsafe prototypes"—the core keywords of a harmful request.[34] However, adversarial prompts strategically divert the model's attention away from these sensitive words.[36]

- **Attention Intensity on Sensitive Words (AttnSensWords)**: This metric measures the weight assigned to tokens like "bomb" or "hack".[36]
- **Attention Dispersion Entropy (AttnEntropy)**: In successful jailbreaks, the attention distribution becomes more dispersed across the benign context or dialogue history, reducing the "signal" of the harmful intent.[36]
- **Mechanism**: During a multi-turn attack, the model gradually reduces the attention it allocates to unsafe requests as the surrounding context (e.g., role-play instructions) expands. This "slipping" prevents the model from triggering its refusal circuitry.[34]

Prefix Forcing and Token Prediction Exploits

Another mechanistic exploit targets the model's causal attention and next-token prediction logic. By forcing the model to start its response with an affirmative phrase like "Sure, I can help with that," the attacker hijacks the reasoning trajectory.[38] Because each subsequent token is conditioned on all previous tokens, the presence of an affirmative prefix makes the generation of a refusal statistically much less likely.[39]

Linguistic Style as an Attack Vector

In 2024 and 2025, research has highlighted that surface-level linguistic styles—rather than just semantic content—can reframe harmful intent and bypass safety guardrails. These "style-augmented" attacks exploit the model's training to be empathetic and helpful.[11]

Linguistic Style	Effectiveness Mechanism	Impact on Success Rate (ASR)
Fearful/Anxious	Mimics distress, causing the model to prioritize "helping" the user over safety policies.[11]	Increases by up to +57 percentage points.[41]
Curious/Inquisitive	Masks harmful intent as a legitimate educational	Significant increase across all benchmarks.[41]

	inquiry.[11]	
Kind/Compassionate	Uses socially pleasant discourse to bypass toxicity triggers.[41]	High effectiveness in naturalistic rewrites.[41]
Pleading/Desperate	Leverages emotional pressure to coerce the model into compliance.[41]	Consistent performance across open-source models.[41]

These attacks are particularly dangerous because they mirror natural human communication, making them difficult for keyword-based or static filters to detect. Contextualized rewrites—where the style is woven into the narrative—consistently outperform simple templates.[41]

Defensive Architectures: Constitutional AI and RLAIF

As adversarial methods become more sophisticated, the focus has shifted toward robust alignment frameworks that bake safety directly into the model's weights. Anthropic's "Constitutional AI" (CAI) represents the gold standard for this approach, utilizing a Reinforcement Learning from AI Feedback (RLAIF) pipeline.[40]

The CAI Training Loop: Critique and Revision

CAI aligns models according to a "constitution"—a set of ethical principles—without requiring extensive human labeling.[46]

1. **Supervised Fine-Tuning (SFT) Phase**:
    - **Initial Response**: The model generates a response to a prompt designed to elicit harmful behavior.[40]
    - **Critique**: The model is asked to critique its own response based on a constitutional principle (e.g., "Discuss whether this response is insensitive or socially inappropriate").[40]
    - **Revision**: The model rewrites the response to align with the principle.[40]
    - **Training**: The base model is fine-tuned on these self-revised "gold standard" examples.[40]
2. **RLAIF Phase**:
    - **Preference Generation**: A "judge" model evaluates pairs of responses based on constitutional adherence.[45]
    - **Reward Model Training**: The AI-generated labels are used to train a safety reward model.[45]
    - **Policy Optimization**: The model is optimized using Proximal Policy Optimization (PPO) or Direct Preference Optimization (DPO) to maximize the reward signal.[45]

Interpretability via QA-LIGN

A newer refinement, QA-LIGN, replaces opaque reward models with structured natural language programs. By operationalizing each constitutional principle into distinct dimensions and evaluative queries, researchers can monitor exactly *why* a model is refusing or accepting a prompt.[50] This transparency improves alignment effectiveness and reduces the "false refusal" rate where harmless queries are incorrectly blocked.[50]

Multi-Layered Security and Guardrail Systems

Beyond model alignment, the industry has adopted a "defense-in-depth" approach involving external guardrails that pre-process inputs and post-process outputs.[51]

Safety Classifiers and Taxonomies

Defense System	Architecture Details	Key Mitigation Goal
Llama Guard	Instruction-tuned classifier based on a customizable risk taxonomy.[54]	Moderating inputs and outputs in real-time pipelines.[54]
Constitutional Classifiers	Input/output filters trained on 10,000 synthetic jailbreak prompts.[44]	Filtering the majority of universal jailbreaks with low overhead.[44]
Sentra-Guard	Hybrid architecture with FAISS-indexed SBERT embeddings and transformer classifiers.[57]	Identifying direct and obfuscated attack vectors with 99.96% detection rate.[57]
JBSHIELD	Detection (D) and Mitigation (M) modules targeting hidden representations.[58]	Strengthening toxic concepts and weakening jailbreak activations in a single pass.[58]

Mechanistic Interventions: Attention Sharpening

Directly countering the "Attention Slipping" phenomenon, "Attention Sharpening" is an inference-time defense that uses temperature scaling to focus the model's attention.[34] By sharpening the attention score distribution, the model is forced to process the unsafe components of a query that it might otherwise "slip" past.[35] This defense is computationally

efficient as it does not require additional model parameters or complex rerouting.[35]

Advanced 2025 Attack Vectors: Reasoning and Multimodal Models

The frontier of jailbreaking research is currently exploring Large Reasoning Models (LRMs) and Multimodal LLMs (MLLMs), which introduce new vulnerabilities through their complex processing stages.

Reasoning Hijacking (H-CoT)

Models that use explicit "Chain-of-Thought" (CoT) reasoning—such as DeepSeek-R1 or o4-mini—can be subverted through "H-CoT" (Fabricated Chain-of-Thought) attacks.[14]

- **Mechanism**: The attacker inserts fabricated reasoning steps into the prompt, disguised as the model's own plausible thought chain.[59]
- **Result**: This allows the attack to bypass initial safety gates by forming a "rational" path to a harmful conclusion before the model produces its final response.[59]

Multimodal Jailbreaking (AERO and JOOD)

As models integrate vision and language, attackers are using images as "carriers" for malicious payloads.[60]

- **FigStep**: Renders malicious text as typography within images and queries the model with neutral prompts like "What does this image show?".[60]
- **AERO (Attention-Enhancement and Targeted Entropy Regularization)**: Specifically models cross-modal attention dynamics to redirect the model's focus toward visual tokens and away from the safety-aligned textual features.[60]
- **Universal Visual Jailbreaks**: A single visual adversarial example can be crafted that, when used as a prefix, compels the model to fulfill a wide range of harmful text instructions that it would otherwise refuse.[61]

Conclusion and Future Outlook for Prevention

The evolution of adversarial prompting from simple "ignore previous instructions" commands to mathematically optimized gradient attacks and psychologically grounded multi-turn steering demonstrates that LLM security is a dynamic and adversarial domain. For university research and future prevention, several key insights emerge:

- **Structural Discrepancy**: Jailbreaking exploits the gap between what a model *can* do and what it is *willing* to do. Narrowing this gap requires embedding alignment rules into the pre-trained weights rather than relying solely on superficial output filters.[3]
- **Mechanism-Specific Defenses**: Understanding internal phenomena like Attention

Slipping allows for the development of targeted, low-overhead defenses like Attention Sharpening that address the root cause of model compliance.[34]

- **Dynamic Scaling**: The power law relationship in Many-Shot Jailbreaking suggests that as models become more capable and handle longer contexts, they become *more* vulnerable to certain classes of attacks, requiring safety measures to scale alongside capabilities.[21]
- **Human-Centric Persuasion**: The success of linguistic style-augmented jailbreaks indicates that alignment must account for human-like communication patterns, such as empathy toward fear or curiosity, which currently serve as blind spots in safety training.[11]

The goal of future security frameworks must be to move toward "interpretable AI safety," where the reasoning behind a model's behavior is transparent and controllable, enabling the development of robust, context-aware guardrails that can withstand the increasingly diverse landscape of adversarial prompt engineering.[50]

Obras citadas

1. LLM Jailbreaking Taxonomy - Innodata, fecha de acceso: enero 25, 2026, https://innodata.com/llm-jailbreaking-taxonomy/
2. A comprehensive analysis of jailbreak vulnerabilities in large language models, fecha de acceso: enero 25, 2026, https://journalijsra.com/sites/default/files/fulltext_pdf/IJSRA-2025-2588.pdf
3. Crescendo attack, fecha de acceso: enero 25, 2026, https://crescendo-the-multiturn-jailbreak.github.io/
4. A Review of “Do Anything Now” Jailbreak Attacks in Large Language Models: Potential Risks, Impacts, and Defense Strategies - Preprints.org, fecha de acceso: enero 25, 2026, https://www.preprints.org/manuscript/202509.0081/v1
5. A Hybrid Perplexity-MAS Framework for Proactive Jailbreak Attack Detection in Large Language Models - MDPI, fecha de acceso: enero 25, 2026, https://www.mdpi.com/2076-3417/15/24/13190
6. Guarding the Guardrails: A Taxonomy-Driven Approach to Jailbreak Detection - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2510.13893v1
7. (PDF) Jailbreak -Unveiling Security Vulnerabilities of ChatGPT - ResearchGate, fecha de acceso: enero 25, 2026, https://www.researchgate.net/publication/394937187_Jailbreak_-Unveiling_Security_Vulnerabilities_of_ChatGPT
8. Evading LLMs' Safety Boundary with Adaptive Role-Play Jailbreaking - MDPI, fecha de acceso: enero 25, 2026, https://www.mdpi.com/2079-9292/14/24/4808
9. What is AI Jailbreaking? - Ironscales, fecha de acceso: enero 25, 2026, https://ironscales.com/glossary/what-is-ai-jailbreaking
10. arXiv:2401.06373v2 [cs.CL] 23 Jan 2024, fecha de acceso: enero 25, 2026, https://par.nsf.gov/servlets/purl/10615999
11. The Fragility of Language: Why LLM Guardrails Are Inherently Vulnerable to Clever Phrasing - Vinova SG, fecha de acceso: enero 25, 2026, https://vinova.sg/the-fragility-of-language-why-llm-guardrails-are-inherently-vul

nerable-to-clever-phrasing/
12. (PDF) A Review of “Do Anything Now” Jailbreak Attacks in Large Language Models: Potential Risks, Impacts, and Defense Strategies - ResearchGate, fecha de acceso: enero 25, 2026, https://www.researchgate.net/publication/395135247_A_Review_of_Do_Anything_Now_Jailbreak_Attacks_in_Large_Language_Models_Potential_Risks_Impacts_and_Defense_Strategies
13. What is the ChatGPT DAN Prompt? Explore 7 Jailbreak Prompts In 2024 - IntellicoWorks, fecha de acceso: enero 25, 2026, https://intellicoworks.com/chatgpt-dan-prompt/
14. Jailbreaking Commercial Black-Box LLMs with Explicitly Harmful Prompts - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2508.10390v3
15. AiActivity/All-Prompt-Jailbreak · Datasets at Hugging Face, fecha de acceso: enero 25, 2026, https://huggingface.co/datasets/AiActivity/All-Prompt-Jailbreak
16. ChatGPT-Dan-Jailbreak.md · GitHub, fecha de acceso: enero 25, 2026, https://gist.github.com/coolaj86/6f4f7b30129b0251f61fa7baaa881516?permalink_comment_id=5579792
17. Great, Now Write an Article About That: The Crescendo Multi-Turn LLM Jailbreak Attack - USENIX, fecha de acceso: enero 25, 2026, https://www.usenix.org/system/files/conference/usenixsecurity25/sec25cycle1-prepub-805-russinovich.pdf
18. A Quiet Crescendo in the Arms Race of LLM Jailbreaking - USENIX, fecha de acceso: enero 25, 2026, https://www.usenix.org/publications/loginonline/crescendo-quiet-crescendo-arms-race-llm-jailbreaking
19. Foot-In-The-Door: A Multi-turn Jailbreak for LLMs - ACL Anthology, fecha de acceso: enero 25, 2026, https://aclanthology.org/2025.emnlp-main.100.pdf
20. 'Echo Chamber' is the Latest Method for Jailbreaking LLMs - Techstrong.ai, fecha de acceso: enero 25, 2026, https://techstrong.ai/prompt-engineering/echo-chamber-is-the-latest-method-for-jailbreaking-llms/
21. Many-shot Jailbreaking - OpenReview, fecha de acceso: enero 25, 2026, https://openreview.net/forum?id=cw5mgd71jW
22. Many-shot jailbreaking \ Anthropic, fecha de acceso: enero 25, 2026, https://www.anthropic.com/research/many-shot-jailbreaking
23. LLM Guardrails Fall to a Simple "Many-Shot Jailbreaking" Attack, Anthropic Warns, fecha de acceso: enero 25, 2026, https://www.hackster.io/news/llm-guardrails-fall-to-a-simple-many-shot-jailbreaking-attack-anthropic-warns-f6eb7b37f4cc
24. 2.4: Scaling Laws | AI Safety, Ethics, and Society Textbook, fecha de acceso: enero 25, 2026, https://www.aisafetybook.com/textbook/scaling-laws
25. How Do Large Language Monkeys Get Their Power (Laws)? - ResearchGate, fecha de acceso: enero 25, 2026, https://www.researchgate.net/publication/389351809_How_Do_Large_Language_Monkeys_Get_Their_Power_Laws

26. How Do Large Language Monkeys Get Their Power (Laws)? - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/pdf/2502.17578
27. How Do Large Language Monkeys Get Their Power (Laws)? - OpenReview, fecha de acceso: enero 25, 2026, https://openreview.net/pdf?id=QqVZ28qems
28. Adversarial Attack on Large Language Models using Exponentiated Gradient Descent, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2505.09820v1
29. The Resurgence of GCG Adversarial Attacks on Large Language Models - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2509.00391v1
30. Greedy Coordinate Gradient Strategy - GCG - Promptfoo, fecha de acceso: enero 25, 2026, https://www.promptfoo.dev/docs/red-team/strategies/gcg/
31. GCG: Adversarial Attacks on Large Language Models | by Brian Pulfer | Medium, fecha de acceso: enero 25, 2026, https://medium.com/@brianpulfer/gcg-adversarial-attacks-on-large-language-models-61f8b51734e9
32. Adversarial Attack on Large Language Models using Exponentiated Gradient Descent, fecha de acceso: enero 25, 2026, https://www.alphaxiv.org/overview/2505.09820v1
33. Adversarial Attacks On LLMs - Lil'Log | PDF | Applied Mathematics - Scribd, fecha de acceso: enero 25, 2026, https://www.scribd.com/document/687957378/Adversarial-Attacks-on-LLMs-Lil-Log
34. Attention Slipping: A Mechanistic Understanding of Jailbreak Attacks and Defenses in LLMs, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2507.04365v1
35. Attention Slipping: A Mechanistic Understanding of Jailbreak Attacks and Defenses in LLMs, fecha de acceso: enero 25, 2026, https://www.researchgate.net/publication/393478071_Attention_Slipping_A_Mechanistic_Understanding_of_Jailbreak_Attacks_and_Defenses_in_LLMs
36. Attention-Based LLM Jailbreak | LLM Security Database - Promptfoo, fecha de acceso: enero 25, 2026, https://www.promptfoo.dev/lm-security-db/vuln/attention-based-llm-jailbreak-6bf6a966
37. Multi-Turn Jailbreaking Large Language Models via Attention Shifting, fecha de acceso: enero 25, 2026, https://ojs.aaai.org/index.php/AAAI/article/view/34553/36708
38. Breaking to Build: A Threat Model of Prompt-Based Attacks for Securing LLMs - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2509.04615v1
39. From Rejection to Acceptance: Model Editing Guided by Representation Transition for Jailbreak Backdooring LLMs | OpenReview, fecha de acceso: enero 25, 2026, https://openreview.net/forum?id=fYWssoFBo6
40. Inside Claude: How Anthropic Built the World's Most Safety-Conscious AI | by NJ - Medium, fecha de acceso: enero 25, 2026, https://medium.com/@nraman.n6/inside-claude-how-anthropic-built-the-worlds-most-safety-conscious-ai-c73f955f9f65
41. Say It Differently: Linguistic Styles as Jailbreak Vectors - arXiv, fecha de acceso:

enero 25, 2026, https://arxiv.org/html/2511.10519v1
42. Say It Differently: Linguistic Styles as Jailbreak Vectors - ResearchGate, fecha de acceso: enero 25, 2026, https://www.researchgate.net/publication/397595848_Say_It_Differently_Linguistic_Styles_as_Jailbreak_Vectors
43. Say It Differently: Linguistic Styles as Jailbreak Vectors - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/pdf/2511.10519
44. Constitutional Classifiers: Defending against universal jailbreaks - Anthropic, fecha de acceso: enero 25, 2026, https://www.anthropic.com/research/constitutional-classifiers
45. RLAIF vs RLHF: What's the Difference and Why It Matters | by PrajnaAI | Medium, fecha de acceso: enero 25, 2026, https://prajnaaiwisdom.medium.com/rlaif-vs-rlhf-whats-the-difference-and-why-it-matters-8e515aa374b6?source=rss------reinforcement_learning-5
46. C3AI: Crafting and Evaluating Constitutions for Constitutional AI - Daniele Quercia, fecha de acceso: enero 25, 2026, https://researchswinger.org/publications/c3ai25.pdf
47. Constitutional AI with Open LLMs - Hugging Face, fecha de acceso: enero 25, 2026, https://huggingface.co/blog/constitutional_ai
48. RLAIF: Scaling Reinforcement Learning from Human Feedback with AI Feedback - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2309.00267v2
49. Comprehensive Guide to Reinforcement Learning in Modern AI - Hugging Face, fecha de acceso: enero 25, 2026, https://huggingface.co/blog/ProCreations/guide-to-rl
50. QA-LIGN: Aligning LLMs through Constitutionally Decomposed QA - ACL Anthology, fecha de acceso: enero 25, 2026, https://aclanthology.org/2025.findings-emnlp.1123.pdf
51. Mitigate jailbreaks and prompt injections - Claude API Docs, fecha de acceso: enero 25, 2026, https://platform.claude.com/docs/en/test-and-evaluate/strengthen-guardrails/mitigate-jailbreaks
52. Model Guardrails Are Getting Better. That Doesn't Mean Your Product Is Safe. - Testlio, fecha de acceso: enero 25, 2026, https://testlio.com/blog/model-guardrails-and-testing/
53. Essential Guide to LLM Guardrails: Llama Guard, NeMo.. | by Sunil Rao - Medium, fecha de acceso: enero 25, 2026, https://medium.com/data-science-collective/essential-guide-to-llm-guardrails-llama-guard-nemo-d16ebb7cbe82
54. LLaMA Guard: Open-Source Safety Classifier - Emergent Mind, fecha de acceso: enero 25, 2026, https://www.emergentmind.com/topics/llama-guard
55. Optimizing Safety Alignment and Jailbreak Defense for Large Language Models, fecha de acceso: enero 25, 2026, https://www.preprints.org/manuscript/202511.1521/v1
56. Anthropic tests algorithm to reliably prevent "AI jailbreaks" - All-About-Industries, fecha de acceso: enero 25, 2026,

https://www.all-about-industries.com/ai-tests-algorithm-to-reliably-prevent-ai-jailbreaks-a-e02b7a592eb0580012750e511b72b776/
57. Sentra-Guard: A Multilingual Human-AI Framework for Real-Time Defense Against Adversarial LLM Jailbreaks - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2510.22628v1
58. JBShield: Defending Large Language Models from Jailbreak Attacks through Activated Concept Analysis and Manipulation - USENIX, fecha de acceso: enero 25, 2026, https://www.usenix.org/system/files/conference/usenixsecurity25/sec25cycle1-prepub-341-zhang-shenyi.pdf
59. AGA: ATTENTION-GUIDED JAILBREAK ATTACKS ON LARGE REASONING MODELS - OpenReview, fecha de acceso: enero 25, 2026, https://openreview.net/pdf/b42fe9d7a0b958c9db7c985ff2525cb007f79afa.pdf
60. Jailbreaking MLLMs via Attention Redirection and Entropy Regularization - MDPI, fecha de acceso: enero 25, 2026, https://www.mdpi.com/2079-9292/15/1/237
61. visual adversarial examples jailbreak aligned large language models - AWS, fecha de acceso: enero 25, 2026, https://moonlight-paper.s3.ap-northeast-2.amazonaws.com/anonymous/2c6d93b3-da76-40df-b1c4-de5d656026ac.pdf
62. Look Before You Leap: Enhancing Attention and Vigilance Regarding Harmful Content with GuidelineLLM - arXiv, fecha de acceso: enero 25, 2026, https://arxiv.org/html/2412.10423v1